

CSI 508. Database Systems I – Spring 2017

Programming Assignment II

The total grade for this assignment is 100 points. The deadline for this assignment is **11:59 PM, April 17, 2017**. *Submissions after this deadline will not be accepted.* Students are required to enter the UAlbany Blackboard system and then upload a .zip file (in the form of [first name]_[last name].zip) that contains the Eclipse project directory and a short document describing:

- any missing or incomplete elements of the code
- any changes made to the original API
- the amount of time spent for this assignment
- suggestions or comments if any

In this programming assignment, you need to complete a Java program that demonstrates how B+-trees run. For example, Figures 1 and 2 show how a B+-tree can change as a key ‘b’ is inserted into the tree.

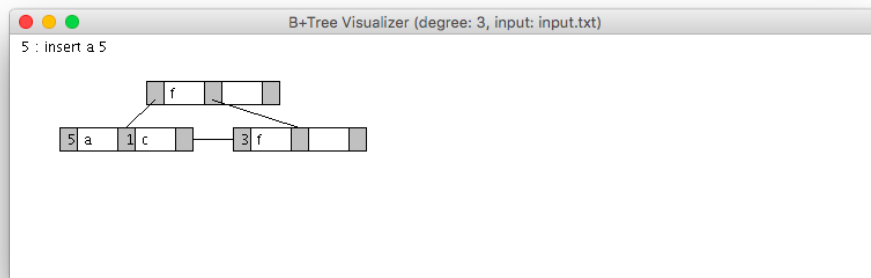


Figure 1: Before Inserting b

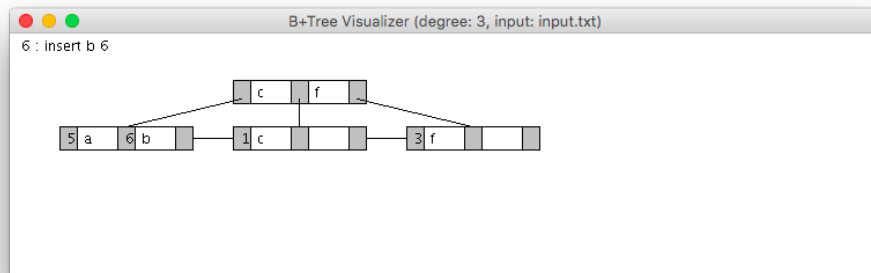


Figure 2: After Inserting b

You first need to run Eclipse on your machine and import the “bplus_tree” project (see Appendix A). This assignment provides an eclipse project that contains a B+-tree visualizer (see

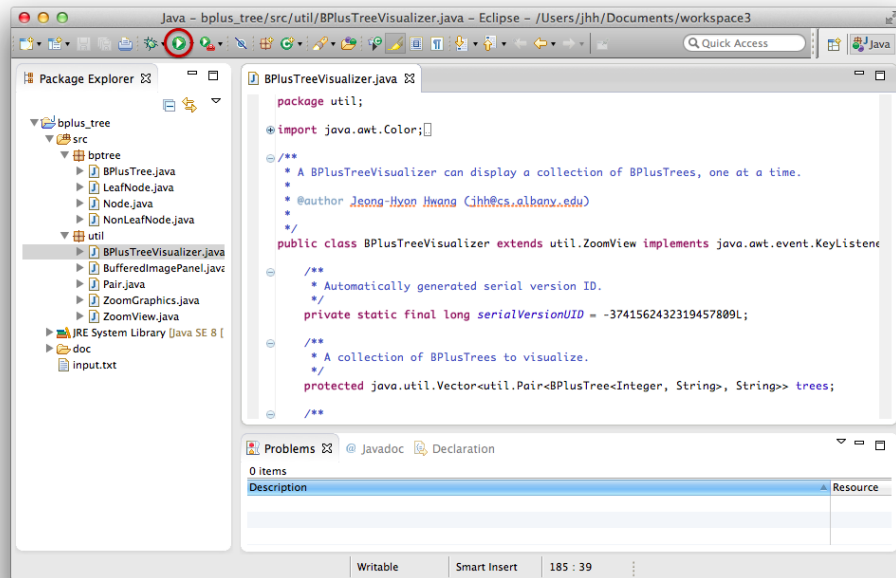


Figure 3: Eclipse

BPlusTreeVisualizer.java) and a partially implemented B+-tree class (see BPlusTree.java). In BPlusTree.java, methods for finding/inserting a key-pointer pair are already implemented (see insert(K k, P p) as well as Algorithms 1, 2, and 3 in Appendix B). For this assignment, you need to implement the delete(K k, P p) method so that we can also remove key-pointer pairs from the tree (refer to Algorithms 4, 5, and 6 in Appendix C). Within your code, please add *detailed comments* so that every step of deletion can be clearly understood. Insufficient comments will negatively affect your grade.

As a side note, the visualizer executes both *insert* and *delete* commands defined in the input.txt file (see the “bplus_tree” project directory). Please ensure that your code works well even if you change the *degree* of the B+-tree and the *insert* and *delete* commands in input.txt.

You can play with the visualizer as follows:

1. move to the previous/next frame (left and right arrow keys, respectively)
2. zoom in/out (up and down arrow keys, respectively, or left and right double clicks)
3. panning ([ctrl] key + left/right/up/down arrow keys, or mouse drag and drop without pressing any key)

Good luck! I hope this assignment will help you better understand how B+-trees run.

Appendix A. Importing the bplus_tree Project

1. Run Eclipse. In the menu bar, choose “File” and then “Import”. Next, select “General” and “Existing Projects into Workspace”. Then, click the “Browse” button and select the “bplus_tree.zip” file contained in this assignment package.
2. Once the project is imported, you can choose BPlusTreeVisualizer.java and then run the program by clicking the icon highlighted in Figure 3.

Appendix B. Pseudo-code for Inserting an Entry in a B+-tree

Algorithm 1: insert(*key* K , *pointer* P)

```

1 if tree is empty then
2   | Create an empty leaf node  $L$ , which is also the root
3 else
4   | Find the leaf node  $L$  that should contain key  $K$ 
5 if  $L$  has less than  $n - 1$  keys then
6   | insert_in_leaf( $L$ ,  $K$ ,  $P$ )
7 else
8   | Copy  $L.P_0 \cdots L.K_{n-2}$  to a block of memory  $T$  that can hold  $n$  (pointer, key) pairs
9   | insert_in_leaf( $T$ ,  $K$ ,  $P$ )
10  | Create node  $L'$ 
11  | Set  $L'.P_{n-1} = L.P_{n-1}$ 
12  | Erase  $L.P_0$  through  $L.K_{n-2}$  from  $L$ 
13  | Set  $L.P_{n-1} = L'$ 
14  | Copy  $T.P_0$  through  $T.K_{\lceil n/2 \rceil - 1}$  from  $T$  into  $L$  starting at  $L.P_0$ 
15  | Copy  $T.P_{\lceil n/2 \rceil}$  through  $T.K_{n-1}$  from  $T$  into  $L'$  starting at  $L'.P_0$ 
16  | Let  $K'$  be the smallest key in  $L'$ 
17  | insert_in_parent( $L$ ,  $K'$ ,  $L'$ )

```

Algorithm 2: insert_in_leaf(*node* L , *key* K , *pointer* P)

```

1 if  $K < L.K_0$  then
2   | Insert  $P$ ,  $K$  into  $L$  just before  $L.P_0$ 
3 else
4   | Let  $K_i$  be the highest key in  $L$  that is less than  $K$ 
5   | Insert  $P$ ,  $K$  into  $L$  just after  $T.K_i$ 

```

Algorithm 3: insert_in_parent(*node* N , *key* K , *node* N')

```

1 if  $N$  is the root of the tree then
2   | Create a new node  $R$  containing  $N$ ,  $K$ ,  $N'$ 
3   | Make  $R$  the root of the tree
4   | return
5 Let  $P = \text{parent}(N)$ 
6 if  $P$  has less than  $n$  pointers then
7   | Insert( $K$ ,  $N'$ ) in  $P$  just after  $N$ 
8 else
9   | Copy  $P$  to a block of memory  $T$  that can hold  $P$  and  $(K, N')$ 
10  | Insert  $(K, N')$  into  $T$  just after  $N$ 
11  | Erase all entries from  $P$ 
12  | Create node  $P'$ 
13  | Copy  $T.P_0 \cdots T.P_{\lceil n/2 \rceil - 1}$  into  $P$ 
14  | Copy  $T.P_{\lceil n/2 \rceil} \cdots T.P_n$  into  $P'$ 
15  | Let  $K' = T.K_{\lceil n/2 \rceil - 1}$ 
16  | insert_in_parent( $P$ ,  $K'$ ,  $P'$ )

```

Appendix C. Pseudo-code for Deleting an Entry in a B+-tree

Algorithm 4: delete(*key* K , *pointer* P)

```
1 find the leaf node  $L$  that contains  $(K, P)$ 
2 delete_entry( $L, K, P$ )
```

Algorithm 5: delete_entry(*node* N , *key* K , *pointer* P)

```
1 delete  $(K, P)$  from  $N$ 
2 if  $N$  is the root then
3   if  $N$  has only one remaining child then
4     Make the child of  $N$  the new root of the tree and delete  $N$ 
5 else if  $N$  has too few keys/pointers then
6   Let  $N'$  be the previous or next child of  $\text{parent}(N)$ 
7   Let  $K'$  be the key between pointers  $N$  and  $N'$  in  $\text{parent}(N)$ 
8   if entries in  $N$  and  $N'$  can fit in a single node then
9     if  $N'$  is the predecessor of  $N$  then
10      merge( $N', K', N$ )
11    else
12      merge( $N, K', N'$ )
13 else
14   // redistribution: move a key and a pointer from node  $N'$  to node  $N$ 
15   if  $N'$  is the predecessor of  $N$  then
16     if  $N$  is a non-leaf node then
17       Let  $m$  be such that  $N'.P_m$  is the last pointer in  $N'$ 
18       Insert  $(N'.P_m, K')$  as the first pointer and key in  $N$  by shifting other pointers and
19       keys right
20       Remove  $(N'.K_{m-1}, N'.P_m)$  from  $N'$ 
21       Replace  $K'$  in  $\text{parent}(N)$  by  $N'.K_{m-1}$ 
22     else
23       Let  $m$  be such that  $(N'.P_m, N'.K_m)$  is the last pointer/key pair in  $N'$ 
24       Insert  $(N'.P_m, N'.K_m)$  as the first pointer and key in  $N$  by shifting other pointers
25       and keys right
26       Remove  $(N'.P_m, N'.K_m)$  from  $N'$ 
27       Replace  $K'$  in  $\text{parent}(N)$  by  $N'.K_m$ 
28   else
29     ... symmetric to the then case ...
```

Algorithm 6: merge(*node* N' , *key* K' , *node* N)

```
1 if  $N$  is not a leaf node then
2   Append  $K'$  and all pointers and keys in  $N$  to  $N'$ 
3 else
4   Append all  $(K_i, P_i)$  pairs in  $N$  to  $N'$ 
5   Set  $N'.P_{n-1} = N.P_{n-1}$ 
6 delete_entry( $\text{parent}(N), K', N$ )
7 delete node  $N$ 
```
